

Technical Overview – System Architecture

System Architecture

- **Data Layer:**
Synthetic soil, rainfall, NDVI, and fertiliser datasets stored as CSV/GeoJSON in data/.
Derived feature table (panels.csv) used for model training.
- **Model Layer:**
scikit-learn Histogram-Based Gradient Boosting Classifier (HGB).
Features: soil nutrients, NDVI trends, rainfall, fertiliser load, slope/flow proxies.
Target: synthetic risk_label generated by heuristic risk function.
- **API Layer:**
FastAPI REST endpoint /score — serves model predictions and generates Alert Card JSONs.
- **Dashboard Layer:**
Streamlit front-end visualising per-field risk status (Green / Amber / Red) and top feature drivers.

Training & Validation

- Training set: 75%, test set: 25% split.
- Model metrics:
 - AUC = 0.998
 - Accuracy = 0.995
 - Precision / Recall / F1 = 0.667 / 0.667 / 0.667
- Feature importance: nitrogen load > recent rainfall > NDVI average.

Technology Stack

- Python 3.11, Pandas 2.2, NumPy 2.0, scikit-learn 1.5
- FastAPI 0.115, Uvicorn 0.30, Streamlit 1.38
- IDE: VS Code (Windows 10 environment)

File Structure

Agric Project/

|

|— data/ # synthetic dataset package

```
└─ outputs/      # model + metrics + plots
└─ generate_data.py  # Step 1: dataset generator
└─ train_model.py   # Step 2: model training
└─ serve_api.py     # Step 3: alert card API
└─ dashboard.py     # Step 4: Streamlit dashboard
└─ requirements.txt
```